

SCIKIT-LEARN API SUMMARY

AVANI NARVEKAR J042

Linear Regression

LinearRegression fits a linear model with coefficients $w = (w_1, \dots, w_p)$ to minimize the residual sum of squares between the observed targets in the dataset, and the targets predicted by the linear approximation.

PARAMETERS:

- fit_intercept: bool, default=True
- normalize: bool, default=False
- copy_X: bool, default=True
- n_jobs: int, default=None
- positive: bool, default=False

ATTRIBUTES:

- coef_: array of shape (n_features,) or (n_targets, n_features)
- rank_int
- singular_: array of shape (min(X, y),)
- intercept_: float or array of shape (n_targets,)

Ridge

Linear least squares with l2 regularization.

Minimizes the objective function:

$$||y - Xw||^2_2 + \alpha * ||w||^2_2$$

This model solves a regression model where the loss function is the linear least-squares function and regularization is given by the l2-norm. Also known as Ridge Regression or Tikhonov regularization. This estimator has built-in support for multi-variate regression (i.e., when y is a 2d-array of shape (n_samples, n_targets)).

PARAMETERS:

- alpha{float, ndarray of shape (n_targets,)}, default=1.0
- fit_intercept: bool, default=True
- normalize: bool, default=False
- copy_X: bool, default=True
- max_iter: int, default=None
- tol: float, default=1e-3
- solver:{'auto', 'svd', 'cholesky', 'lsqr', 'sparse_cg', 'sag', 'saga'}, default='auto'

- `random_state`:int, RandomState instance, default=None

ATTRIBUTES:

- `coef_`:ndarray of shape (n_features,) or (n_targets, n_features)
- `intercept_`:float or ndarray of shape (n_targets,)
- `n_iter_`:None or ndarray of shape (n_targets,)

Lasso

Linear Model trained with L1 prior as regularizer (aka the Lasso)

The optimization objective for Lasso is:

$$(1 / (2 * n_samples)) * ||y - Xw||^2_2 + \alpha * ||w||_1$$

PARAMETERS:

- `alpha`:float, default=1.0
- `fit_intercept`:bool, default=True
- `normalize`:bool, default=False
- `precompute`:bool or array-like of shape (n_features, n_features), default=False
- `copy_X`:bool, default=True
- `max_iter`:int, default=1000
- `tol`:float, default=1e-4
- `warm_start`:bool, default=False
- `positive`:bool, default=False
- `random_state`:int, RandomState instance, default=None
- `selection`{'cyclic', 'random'}, default='cyclic'

ATTRIBUTES:

- `coef_`:ndarray of shape (n_features,) or (n_targets, n_features)
- `dual_gap_`:float or ndarray of shape (n_targets,)
- `sparse_coef_`:sparse matrix of shape (n_features, 1) or (n_targets, n_features)
- `intercept_`:float or ndarray of shape (n_targets,)
- `n_iter_`:int or list of int

Logistic Regression:

Logistic Regression (aka logit, MaxEnt) classifier.

In the multiclass case, the training algorithm uses the one-vs-rest (OvR) scheme if the 'multi_class' option is set to 'ovr', and uses the cross-entropy loss if the 'multi_class' option is set to 'multinomial'. (Currently the 'multinomial' option is supported only by the 'lbfgs', 'sag', 'saga' and 'newton-cg' solvers.)

PARAMETERS:

- `Penalty` : {'l1', 'l2', 'elasticnet', 'none'}, default='l2'
- `dual` : bool, default=False
- `tol` : float, default=1e-4
- `C` : float, default=1.0
- `fit_intercept` : bool, default=True
- `intercept_scaling` : float, default=1
- `class_weightdict` or 'balanced', default=None
- `random_state` : int, RandomState instance, default=None
- `solver`{'newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga'}, default='lbfgs'
- `max_iter` : int, default=100
- `multi_class`{'auto', 'ovr', 'multinomial'}, default='auto'
- `verbose` : int, default=0
- `warm_start` : bool, default=False
- `n_jobs` : int, default=None
- `l1_ratio` : float, default=None

ATTRIBUTES:

- `classes_` ndarray of shape (n_classes,)
- `coef_` ndarray of shape (1, n_features) or (n_classes, n_features)
- `intercept_` ndarray of shape (1,) or (n_classes,)
- `n_iter_` ndarray of shape (n_classes,) or (1,)